

Tame Your Data With R: ADVANCE Workshop 2

Elizabeth E. Esterly

December 4th, 2017, University of New Mexico

```
library(tidyverse)

## Loading tidyverse: ggplot2
## Loading tidyverse: tibble
## Loading tidyverse: tidyr
## Loading tidyverse: readr
## Loading tidyverse: purrr
## Loading tidyverse: dplyr

## Conflicts with tidy packages -----

## filter(): dplyr, stats
## lag():      dplyr, stats

super <- read_csv('superheroes.csv')

## Parsed with column specification:
## cols(
##   name = col_character(),
##   alignment = col_character(),
##   sex = col_character(),
##   age = col_integer(),
##   publisher = col_character()
## )
```

Transform

We'll bring in a second dataset here, transform it, and then transform our original dataset by doing a join on superhero names. For brevity's sake, assume this one is already tidied up.

```
#Import the second set and take a look at the first few rows
second <- read_csv("~/Desktop/characteristics.csv")
second %>% head()
```

```
## # A tibble: 6 x 3
##       name      status measurement
##   <chr>      <chr>      <chr>
## 1 Wonder Woman homePlanet Themyscira
## 2 Thor        homePlanet Asgard
## 3 Storm       homePlanet Earth
## 4 Superman    homePlanet Krypton
## 5 Deadpool    homePlanet Earth
## 6 Wonder Woman flyingSpeedmph 5,100,000
```

Huh. It looks like the name column repeats itself. Let's get a wider view...

```
second

## # A tibble: 10 x 3
##       name      status measurement
```

```
##           <chr>           <chr>           <chr>
## 1 Wonder Woman   homePlanet Themyscira
## 2      Thor      homePlanet   Asgard
## 3      Storm     homePlanet   Earth
## 4    Superman   homePlanet   Krypton
## 5    Deadpool   homePlanet   Earth
## 6 Wonder Woman flyingSpeedmph 5,100,000
## 7      Thor     flyingSpeedmph 24000
## 8      Storm   flyingSpeedmph 1000
## 9    Superman flyingSpeedmph 7,200,000
## 10   Deadpool  flyingSpeedmph 0
```

spread and gather

It definitely does. This data is a great candidate for transformation. Let's break up that status column into separate factors:

```
#Import the second set and take a look at the first few rows
secondWide <- second %>% spread(status, measurement)
secondWide
```

```
## # A tibble: 5 x 3
##       name flyingSpeedmph homePlanet
## *    <chr>           <chr>           <chr>
## 1   Deadpool             0           Earth
## 2     Storm           1000           Earth
## 3   Superman       7,200,000         Krypton
## 4      Thor           24000         Asgard
## 5 Wonder Woman   5,100,000 Themyscira
```

The Tidyverse `gather` function will help you do the opposite—go from wide to long. To learn more about that function, you can type:

```
help("gather")
```

You can do this with most functions in R to get some more details about them.

Note that the result of `spread` was sorted alphabetically by the first column. We really care about who's the fastest, though:

```
#Sort by speed
#ascending order
secondWide <- secondWide %>% arrange(flyingSpeedmph)
secondWide
```

```
## # A tibble: 5 x 3
##       name flyingSpeedmph homePlanet
##      <chr>           <chr>           <chr>
## 1   Deadpool             0           Earth
## 2     Storm           1000           Earth
## 3      Thor           24000         Asgard
## 4 Wonder Woman   5,100,000 Themyscira
## 5    Superman       7,200,000         Krypton
```

```
#descending order
secondWide <- secondWide %>% arrange(desc(flyingSpeedmph))
secondWide
```

```
## # A tibble: 5 x 3
##       name flyingSpeedmph homePlanet
##       <chr>         <chr>         <chr>
## 1    Superman      7,200,000      Krypton
## 2 Wonder Woman    5,100,000 Themyscira
## 3      Thor        24000      Asgard
## 4      Storm        1000      Earth
## 5    Deadpool          0      Earth
```

This looks like what we want. Let's merge with our `super` dataset.

joins

Don't be scared!

inner join(x, y)

From the documentation:

```
help("inner_join")
```

Return all rows from x where there are matching values in y, and all columns from x and y. If there are multiple matches >between x and y, all combination of the matches are returned. This is a mutating join.

We'll give our `super` dataset as x, because we want to keep all of the characters in `super` that also appear in `secondWide`. `secondWide` will be given second as y.

What is a mutating join? It's some database terminology. It just means you can add variables to the LHS.

```
#Try an inner join
```

```
inner_join(super, secondWide)
```

```
## Joining, by = "name"
```

```
## # A tibble: 4 x 7
##       name alignment  sex  age publisher flyingSpeedmph homePlanet
##       <chr>    <chr> <chr> <int>    <chr>         <chr>         <chr>
## 1    Storm    good female  300    Marvel          1000      Earth
## 2     Thor    good  male   NA    Marvel          24000    Asgard
## 3 Wonder Woman good female 5000      DC      5,100,000 Themyscira
## 4    Superman good  male   29      DC      7,200,000 Krypton
```

```
#The Tidyverse way: use pipes for first function argument 'x'
```

```
#Doesn't make a difference here but a good habit to get into
```

```
super %>% inner_join(secondWide)
```

```
## Joining, by = "name"
```

```
## # A tibble: 4 x 7
##       name alignment  sex  age publisher flyingSpeedmph homePlanet
##       <chr>    <chr> <chr> <int>    <chr>         <chr>         <chr>
## 1    Storm    good female  300    Marvel          1000      Earth
## 2     Thor    good  male   NA    Marvel          24000    Asgard
## 3 Wonder Woman good female 5000      DC      5,100,000 Themyscira
## 4    Superman good  male   29      DC      7,200,000 Krypton
```

Notice we were automatically joined on name. We kept all characters that appeared in *both* datasets and merged all attributes given in both.

semi_join(x, y)

return all rows from x where there are matching values in y, keeping just columns from x

```
super %>% semi_join(secondWide)
```

```
## Joining, by = "name"
```

```
## # A tibble: 4 x 5
```

	name	alignment	sex	age	publisher
	<chr>	<chr>	<chr>	<int>	<chr>
## 1	Storm	good	female	300	Marvel
## 2	Thor	good	male	NA	Marvel
## 3	Wonder Woman	good	female	5000	DC
## 4	Superman	good	male	29	DC

Here, our names were filtered by just the ones that appeared in both datasets, but we didn't pick up the extra columns from the `secondWide` dataset.

anti_join(x, y)

return all rows from x where there are not matching values in y, keeping just columns from x.

```
super %>% anti_join(secondWide)
```

```
## Joining, by = "name"
```

```
## # A tibble: 10 x 5
```

	name	alignment	sex	age	publisher
	<chr>	<chr>	<chr>	<int>	<chr>
## 1	Black Widow	<NA>	female	38	Marvel
## 2	Magneto	bad	male	93	Marvel
## 3	Mystique	bad	female	120	Marvel
## 4	Loki	bad	male	1048	Marvel
## 5	Batman	good	male	32	DC
## 6	Joker	bad	male	46	DC
## 7	0	good	0	NA	DC
## 8	Harley Quinn	bad	female	27	DC
## 9	Supergirl	good	female	24	DC
## 10	Hellboy	good	male	10	Dark Horse

In this case, we exclude all characters that appear in both x and y but still keep just columns from x.

full_join(x, y)

return all rows and all columns from both x and y. Where there are not matching values, returns NA for the one missing.

```
super %>% full_join(secondWide)
```

```
## Joining, by = "name"
```

```
## # A tibble: 15 x 7
```

	name	alignment	sex	age	publisher	flyingSpeedmph
	<chr>	<chr>	<chr>	<int>	<chr>	<chr>
## 1	Black Widow	<NA>	female	38	Marvel	<NA>
## 2	Magneto	bad	male	93	Marvel	<NA>
## 3	Storm	good	female	300	Marvel	1000
## 4	Mystique	bad	female	120	Marvel	<NA>

```
## 5      Thor      good  male    NA      Marvel      24000
## 6      Loki      bad   male  1048     Marvel      <NA>
## 7 Wonder Woman  good female 5000      DC        5,100,000
## 8      Batman    good  male    32      DC          <NA>
## 9      Joker     bad   male    46      DC          <NA>
## 10     0         good    0     NA       DC          <NA>
## 11 Harley Quinn  bad female 27      DC          <NA>
## 12 Supergirl     good female 24      DC          <NA>
## 13 Superman     good  male    29      DC        7,200,000
## 14 Hellboy      good  male    10 Dark Horse <NA>
## 15 Deadpool     <NA> <NA>    NA      <NA>      0
## # ... with 1 more variables: homePlanet <chr>
```

Let's do an anti join one more time and save the result:

```
result <- super %>% anti_join(secondWide)
```

```
## Joining, by = "name"
```

```
result
```

```
## # A tibble: 10 x 5
##       name alignment    sex   age publisher
##       <chr>      <chr> <chr> <int>      <chr>
## 1 Black Widow  <NA> female   38      Marvel
## 2 Magneto      bad   male   93      Marvel
## 3 Mystique     bad female  120      Marvel
## 4 Loki         bad   male  1048     Marvel
## 5 Batman       good  male   32       DC
## 6 Joker        bad   male   46       DC
## 7 0           good    0    NA       DC
## 8 Harley Quinn  bad female  27       DC
## 9 Supergirl     good female  24       DC
## 10 Hellboy     good  male   10 Dark Horse
```

Our dataset has come a long way since the beginning! Let's close by exploring some visualizations.


Visualize

The `ggplot2` library provides an extensive set of tools to help you create beautiful visualizations. It can be hard to get started though, so I suggest looking at great examples, like from this website: <http://r-statistics.co/Top50-Ggplot2-Visualizations-MasterList-R-Code.html#Ordered%20Bar%20Chart> and working backwards from there.

Categorical bar chart

Let's compare alignment between the sexes.

```
result %>% ggplot(aes(sex)) +
  geom_bar(aes(fill=alignment))
```




Lollipop charts

```
result %>% ggplot(aes(x=name, y=age, label=age )) +
  geom_point(stat='identity', fill="black", size=8) +
  geom_segment(aes(y = 0,
                  x = name,
                  yend = age,
                  xend= name),
              color = "black") +
  coord_flip()
```

Warning: Removed 1 rows containing missing values (geom_point).

Warning: Removed 1 rows containing missing values (geom_segment).




Loki is throwing things off here. Let's take him out:

```
result %>% filter(name != "Loki") %>% ggplot(aes(x=name, y=age, label=age )) +
  geom_point(stat='identity', fill="black", size=8) +
  geom_segment(aes(y = 0,
                  x = name,
                  yend = age,
                  xend= name),
              color = "black") +
  coord_flip()
```

```
## Warning: Removed 1 rows containing missing values (geom_point).
```

```
## Warning: Removed 1 rows containing missing values (geom_segment).
```



Much better!

Let's add a little text in the dots for clarity and set our 'baseline' age to 32:

```
result %>% filter(name != "Loki") %>% ggplot(aes(x=name, y=age, label=age )) +
  geom_point(stat='identity', fill="black", size=8) +
  geom_segment(aes(y = 32,
    x = name,
    yend = age,
    xend= name),
    color = "black") +
  geom_text(color="white", size=3)+
  labs(title="Who's older and younger than the Batman?", subtitle="Batman is 32")+
  coord_flip()
```

```
## Warning: Removed 1 rows containing missing values (geom_point).
```

```
## Warning: Removed 1 rows containing missing values (geom_segment).
```

```
## Warning: Removed 1 rows containing missing values (geom_text).
```


Who's older and younger than the Batman?

Batman is 32

